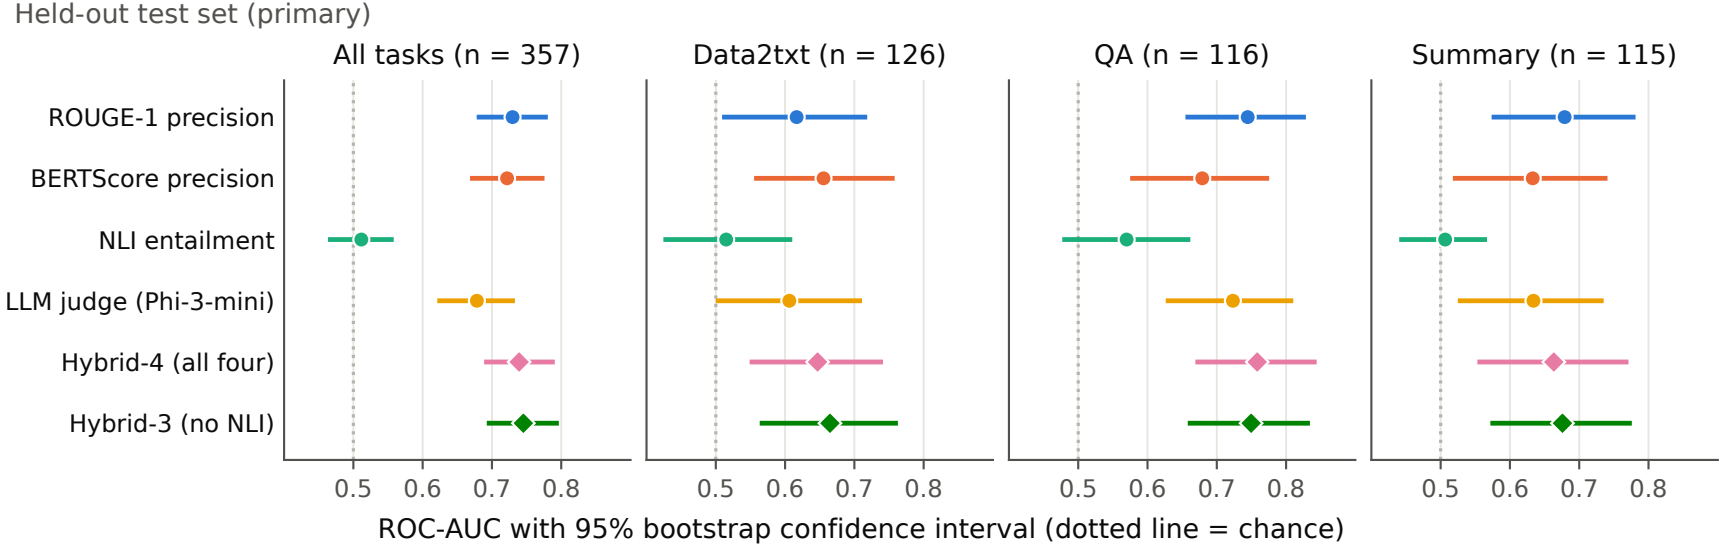
Held-out test set (primary)
All tasks (n = 357)
Data2txt (n = 126)
QA (n = 116)
Summary (n = 115)
ROUGE-1 precision
BERTScore precision
NLI entailment
LLM judge (Phi-3-mini)
Hybrid-4 (all four)
Hybrid-3 (no NLI)
0.5
0.6
0.7
0.8
ROC-AUC with 95% bootstrap confidence interval (dotted line = chance)